POLITECNICO DI TORINO

Master's Degree in Computer Engineering

Master's Degree Thesis

Acceleration by Separate-Process Cache for Memory-Intensive Algorithms on FPGA via High-Level Synthesis

Supervisor
Prof. Luciano Lavagno

Candidate
Giovanni Brignone
ID: 274148

Academic year 2020-2021

Abstract

The end of the Moore's Law validity is making the performance advance of Software run on general purpose processors more challenging than ever. Since current technology cannot scale anymore it is necessary to approach the problem from a different point of view: application-specific Hardware can provide higher performance and lower power consumption, while requiring higher design efforts and higher deployment costs.

The problem of the high design efforts can be mitigated by the High-Level Synthesis (HLS), since it helps improve designer productivity thanks to convenient Software-like tools.

The problem of high deployment costs can be tackled with FPGAs, which allow implementing special-purpose Hardware modules on general-purpose underlying physical architectures.

One of the open issues of HLS is the memory bandwidth bottleneck which limits performance, especially critical in case of memory-bound algorithms.

FPGAs memory system is composed of three main kinds of resources: registers, Block RAMs and external DRAMs. Current HLS tools allow exploiting this memory hierarchy manually, in a scratchpad-like fashion: the objective of this thesis work is to automate the memory management by providing an easily integrable and fully customizable cache system for HLS.

The proposed implementation has been developed using Vitis™HLS tool by Xilinx Inc..

The first development phase produced a single-port cache module, in the form of a C++ class configurable through templates in terms of number of sets, ways, words per line and replacement policy. The cache lines have been mapped to BRAMs. To obtain the desired performance, an unconventional (for HLS) multiprocess architecture has been developed: the cache module is a separate process with respect to the algorithm using it: the algorithm logic sends a memory access request to the cache and reads its response, communicating through FIFOs.

In the second development phase, the focus was put on performance optimization, in two dimensions: increasing the memory hierarchy depth by introducing a Level 1 (L1) cache and increasing parallelism by enabling multiple ports.

The L1 cache is composed of cache logic inlined in the user algorithm: this solution allows to cut the costs of FIFOs communications. To keep L1 cache simple it has been implemented with a write-through write policy, therefore it provides advantages for read accesses only. It is configurable in the number of lines and each line contains the same number of words of the associated Level 2 (L2) cache.

The multi-port solution provides a single L2 cache accessible from multiple FIFOs ports, each of which can be associated with a dedicated L1 cache. It is possible to specify the number of ports through a template parameter and it typically corresponds to the unrolling factor of the loop in which the cache is accessed.

In order to evaluate performance and resource usage impact of the developed cache module, multiple algorithms with different memory access patterns have been synthesized and simulated, with all data accessed to DRAM (performance lower bound), to BRAM (performance higher bound) and to cache (with multiple configurations).

Contents

List of Figures

List of Tables

List of Acronyms

API Application Programming Interface

AXI Advanced eXtensible Interface

BRAM Block RAM

CC Clock Cycle

DRAM Dynamic RAM

DSP Digital Signal Processor

FF Flip-Flop

FIFO First-In First-Out

FPGA Field-Programmable Gate Array

FSM Finite-State Machine

HDL Hardware Description Language

HLS High-Level Synthesis

HW Hardware

II Initiation Interval

IPC Inter-Process Communication

L1 Level 1

L2 Level 2

LRU Least Recently Used

LSB Least Significant Bit

LUT Lookup Table

MSB Most Significant Bit

RAM Random Access Memory

RAW Read After Write

RTL Register-Transfer Level

SW Software

1 Background

The literature about cache systems, the High-Level Synthesis state of the art and an analysis of the resources available on board modern FPGAs are the fundamental background for this thesis work.

1.1 Cache

Memory devices are crucial components of computing systems as they can pose a higher bound in terms of performance, especially when executing memory-intensive algorithms. The ideal memory should be fast, large and cheap, but current technology forces the designer to choose a trade-off between the metrics.

A common solution to this problem is to set up a memory hierarchy in which fast but small memories are paired with large but slow memories, which allows getting good performance on average while containing costs.

This hierarchy can be managed by two main approaches:

- *Scratchpad*: different memories belongs to different addressing spaces: the user is in charge of manually choosing what memory to access: this approach allows to optimally exploit the hierarchy at the cost of high design effort.

- *Cache*: different memories belongs to the same addressing space: the system automatically uses the whole hierarchy, exploiting spatial locality (accessed data is likely physically close to previously accessed data) and temporal locality (accessed data has likely recently been accessed), which are typical of many algorithms.

1.1.1 Structure

A cache memory is logically split into *sets* containing *lines* (or *ways*) which are in turn made up of *words*, as shown in Figure 1.1.

Whenever a word w is requested, there are two possibilities:

- *Hit*: w is present in the cache: the request can be immediately fulfilled.

- *Miss*: w is not present in the cache: it is necessary to retrieve it from lower level memory before fulfilling the request.

Figure 1.1: Cache logic structure.

During the data retrieving, a cache line is filled with a block of contiguous words loaded from the lower level memory, trying to exploit spatial locality of future accesses, while mapping policies and replacement policies determine which cache line to overwrite, trying to exploit temporal locality.

If the cache memory is writable, data consistency is ensured by a consistency policy.

1.1.2   Policies

Mapping policy

The mapping policy is in charge of statically associating a lower level memory line to a cache set.

The *set associative* policy is the most common mapping policy: given a cache memory with s sets of w words, the word address (referred to the lower level memory) bits are split into three parts (as shown in Figure 1.2):

1. $\log_2(w)$: offset of the word in the line.

2. $\log_2(s)$: set.

3. Remaining MSBs: tag identifying the specific line.

Special cases of this policy are:

- *Direct mapped* policy: each set is composed of a single line: the set bits identify a specific cache line, therefore there is no need for a replacement policy.

2

Figure 1.2: Set associative policy address bits meaning.

- *Fully associative* policy: there is only a single set, therefore the line is fully determined by the replacement policy.

Replacement policy

The replacement policy is in charge of dynamically associating a lower level memory line to a cache line of a set.

Multiple solutions of this problem have been developed, trying to maximize the temporal locality exploitation. Among the most commonly used solutions there are:

- *First-In First-Out*: the line to be replaced is the first one that has been inserted to the cache.

- *Least Recently Used*: the line to be replaced is the one that has least recently been accessed.

Consistency policy

The consistency policy is in charge of ensuring data consistency between memories belonging to different hierarchy levels.

The most common solutions to this problem are:

- *Write-back*: write accesses are performed to the highest level memory and lower level memories are updated when the cache line is replaced only.

- *Write-through*: each write access is propagated along the whole hierarchy.

1.1.3 Benefits

A two-level memory hierarchy is composed of a L1 cache memory (access time: t_{L1}; access energy: E_{L1}) and a L2 memory (access time: t_{L2}; access energy: E_{L2}), with $t_{L1} << t_{L2}$ and $E_{L1} << E_{L2}$.

This memory hierarchy is accessed n_{tot} times and n_{hit} of these accesses are cache hits.

The *hit ratio* is defined as:
$$H := \frac{n_{\text{hit}}}{n_{\text{tot}}} \tag{1.1}$$

The *average access time* and *energy* are defined as:

$$\begin{cases} \overline{t}(H) := Ht_{\mathrm{L1}} + (1-H)t_{\mathrm{L2}} \\ \overline{E}(H) := HE_{\mathrm{L1}} + (1-H)E_{\mathrm{L2}} \end{cases} \tag{1.2}$$

Equation 1.2 shows the criticality of the *hit ratio*: the performance and power consumption advantages provided by the cache are significant if and only if H is sufficiently near to 1.

1.2 Field-Programmable Gate Array

Field-Programmable Gate Arrays are integrated circuits able to implement special purpose circuits described in Hardware Description Language (HDL), thanks to their programmable logic blocks and interconnections.

1.2.1 Memory system

A FPGA memory system is typically made up of:

- Registers: the fastest but most expensive memories, therefore they are only a few.

- BRAMs: on chip Random Access Memories (RAMs) accessible through simple and fast interface.

- External DRAMs: off chip DRAMs accessible through complex and slow interface (e.g. AXI).

1.3 High-Level Synthesis

High-Level Synthesis (HLS) is an Electronic Design Automation technique aimed at translating an algorithm description in a high-level Software programming language (such as C and C++) into a HDL description.

HLS allows designing more complex systems in less time, compared to HDL design, moreover makes the Hardware and Software co-design easier, at the cost of limited low-level control.

This Section is mainly referred to *Vitis™ HLS 2020.2* [2] and *2021.1* [3], but most currently available HLS commercial tools provide equivalent features.

1.3.1 Workflow

The typical HLS workflow consists of:

1. *SW implementation*: the top-level entity is a C function: the function arguments are the entity ports and the functionality is implemented in SW; in order to guarantee synthesizability some constraints should be respected (e.g. no dynamic memory allocation).

2. *SW verification*: the testbench can be developed as a simple main function which calls the top-level entity function, therefore the functionality is verified like any SW: it is possible to exploit traditional tools (e.g. debuggers, print statements...).

3. *HW synthesis*: the synthesizer generates a Register-Transfer Level (RTL) description of the top-level entity. It is possible to generate different architectures by setting up some parameters through dedicated directives.

4. *HW verification*: the RTL description is simulated, to make sure that SW and HW outputs match.

1.3.2 Optimization techniques

HLS tools provide different optimization techniques which can be set up by means of compiler directives.

Pipelining

Given a set of sequential stages (e.g. A, B and C of Figure 1.3) which compose an operation (e.g. A + B + C of Figure 1.3) which has to be executed multiple times, the pipelining technique inserts pipeline registers at the output of each stage, so that each stage can run in parallel on different input data (e.g. at the third clock cycle, while C is processing first input, B is processing second input and A is processing third input). The introduced parallelism allows to increase the throughput at a limited additional area cost (only pipeline registers and a FSM are required).

The throughput is determined by the interval (expressed in number of clock cycles) between the beginning of two consecutive executions of the operation, which is called Initiation Interval (II). The optimal pipeline has an II equal to one: at the steady state, one output per clock cycle is produced.

The pipelining can be performed at instruction level, within a loop or a function, or at function level (in HLS terminology this particular kind of pipelining is called *Dataflow*).
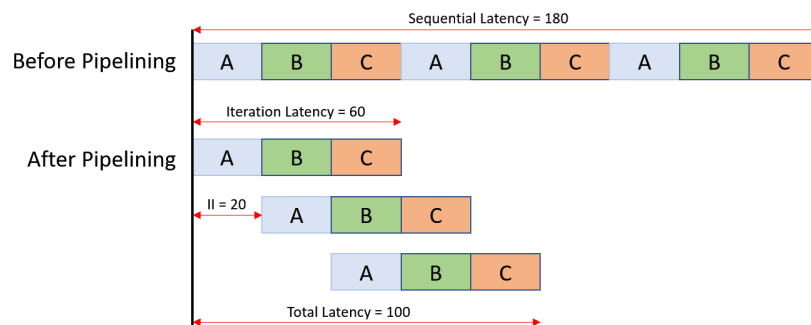
Figure 1.3: Pipelining example.

5

Loop unrolling

The logic of a rolled loop allows the execution of one iteration at a time: if the loop iterates N times and each iteration has a latency L_{it}, the total loop latency is equal to $L_{loop,rolled} := N \cdot L_{it}$.

The loop unrolling technique instantiates the logic for executing f iterations at a time (where f is the unrolling factor). If there are no dependencies between different iterations, the latency of the unrolled loop is: $L_{loop,unrolled}(f) := \frac{N}{f} \cdot L_{it}$.

Loop unrolling can improve both latency and throughput, but it is expensive in terms of resource usage, since they are multiplied by f.
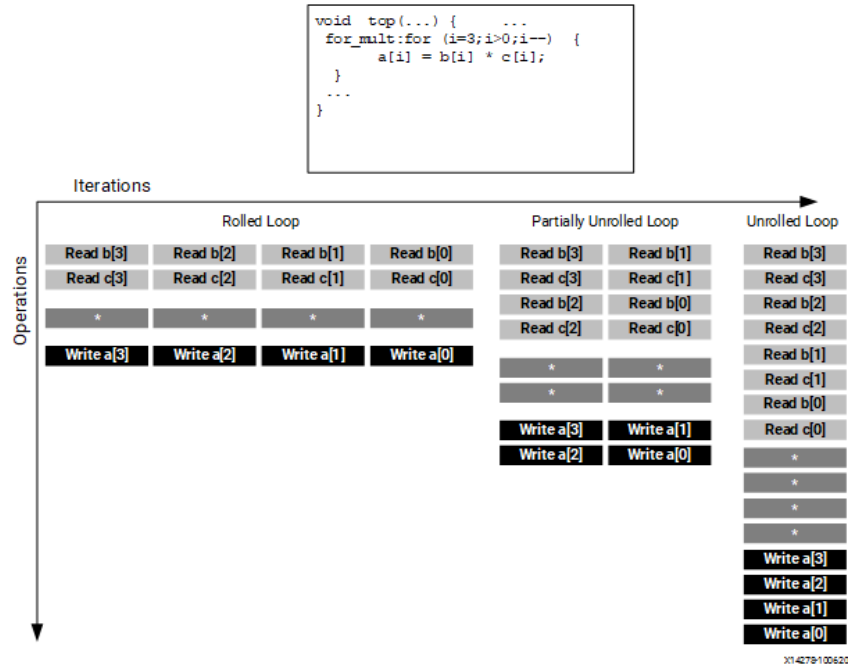
Figure 1.4: Loop unrolling example.

Memory optimizations

- On-chip memory:

    - **Array partitioning**: given a partitioning factor f, an array is split into f portions, each one mapped to a dedicated memory element.

      This allows multiple concurrent accesses to the same array, at the cost of higher memory elements usage.

      Figure 1.5 shows different partitioning modes.

- Off-chip memory:

Figure 1.5: Array partitioning examples.

– **Interface widening**: multiple data elements are packed into a single bigger word, to perform multiple accesses at the same time.

– **Burst accesses**: multiple memory accesses are aggregated into AXI bursts to reduce overall latency and improving throughput.
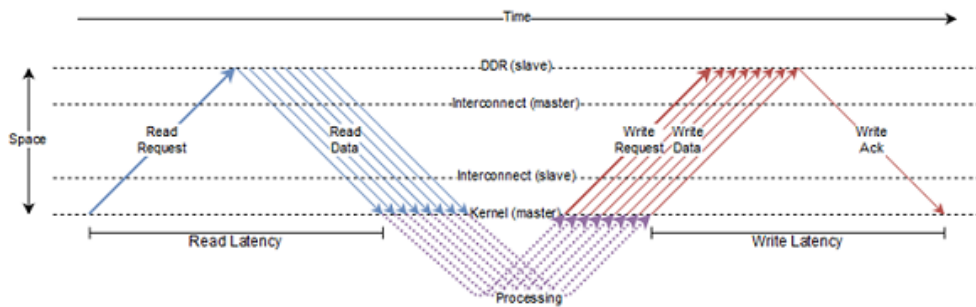
Figure 1.6: Burst read and write example.

2 Motivation

HLS tools are currently unable to automatically exploit the memory hierarchy present on FPGAs: the only way to take advantage of them is the manual management in a *scratchpad*-like manner, which requires additional design and verification efforts.

The proposed solution **automates the low-level memory management** through a cache module for HLS, which works as an interface with the off-chip DRAM (accessible through an AXI bus) and stores its data to on-chip BRAMs and registers.

The proposed cache module has the **dual purpose** of:

- *Reducing the number of DRAM accesses*: misses only needs to access DRAM.

- *Optimizing DRAM accesses*: lines are accessed in bursts through a widened memory interface.

FPGAs provide multiple DRAM ports and HLS can assign each array to a different port: this allows implementing **array-specific** caches, which in general can be easily tuned to reach high hit ratios, since access patterns to a single array are usually regular and there is no interference between accesses to different arrays.

A special attention has been put on **user-friendliness**:

- *Configurability*: cache characteristics can be set through parameters.

- *Integrability*: cache can be inserted into existing designs without requiring many changes.

- *Observability*: critical cache data (e.g. hit ratio) can be profiled during SW simulation for easing the cache parameters tuning.

2.1   Ma's cache

Liang Ma et al. proposed a `C++` cache implementation [6] compatible with *Vivado™HLS 2016.2*.

It is an array-specific cache module in the form of different `C++` classes: each of them implements an access type (read only/write only and read write) and a mapping policy (direct mapped and set associative).

To improve the *integrability* the `operator[]` has been overloaded so that the cache object can be accessed in the same way as array variables, minimizing the required changes to the code which integrates the cache.

This architecture is **inlined**: the cache logic is directly inserted in the user algorithm logic. This is the major limitation of this solution, since the additional logic inserted in the algorithm may make it too complex and worsen the generated circuit performance.

2.2 Proposed solution

The primary goal of this thesis work is to develop the *Basic cache*, a cache architecture which runs in a separate process with respect to the application using it, trying to solve the main limitation of *Ma's cache*: the application logic cluttering due to the inlining.
This architecture has been then optimized in two dimensions:

- *Multi-levels cache*: a L1 cache are added to the cache hierarchy, with the objective of further reducing memory access latency.

- *Multi-ports cache*: multiple cache access points are added to the cache, each one with a dedicated L1 cache, so that multiple requests can be served in parallel.

3 Basic cache

The *Basic cache* is aimed at solving the main limitation of *Ma's cache*: application logic cluttering due to inlining.

3.1 Architecture

The fundamental idea behind the *Basic cache* is that the cache logic is inserted in a separate process with respect to the application logic accessing it (Figure 3.1): this isolation makes the cache always perform in the same manner, independently of the algorithm accessing it, while keeping the application logic as clean as possible, since application only has to write requests to cache and read responses, instead of integrating the whole cache logic.

Figure 3.1: *Single-process Basic cache* architecture.

3.1.1 Functionality

If application A needs to access the array associated with the cache C:

1. *A* sends the access request to *C*: operation (i.e. read or write), address and (in case of write access) data.

2. *C* receives the request and checks if the requested address causes a miss.

3. (in case of miss) *C* prepares its BRAM memory for fulfilling the requested access:

   - (if needed) writes back to DRAM the BRAM line to be replaced.
   - reads from DRAM the requested line and store it to BRAM.

4. *C* performs the requested access to BRAM and (in case of read request) sends requested data to *A*.

3.1.2   Characteristics

The *Basic cache* is compliant with the set associative mapping policy and the write-back consistency policy. It is configurable in terms of:

- Word type and number of words per line.

- Number of sets and ways (therefore, it is possible to obtain a fully associative policy by setting the number of sets to 1 or a direct mapped policy by setting the number of ways to 1).

- Replacement policy (Least Recently Used or First-In First-Out).

3.1.3   Single-process Basic cache

The *Single-process Basic cache* is composed of a single pipelined process which performs all the cache functionalities.

This process can be pipelined with an II equal to 1 when:

- Memory accesses are Read-Only.

- A cache line can fit a single AXI transaction (i.e. line is not bigger than the maximum AXI interface width: 512 or 1024 bits typically, depending on the specific device).

Write accesses generate some dependencies on the AXI interface, while large cache lines require multiple AXI transactions: both of them cause an increase of the cache process II, reducing cache performance.

3.1.4 Multi-processes Basic cache

The *Multi-processes Basic cache* splits cache into two processes (Figure 3.2):

- *Core* process: manages communication with application and keeps cache data structures up to date.

- *Memory interface* process: deals with the AXI interface.

This architecture is aimed at solving the performance limitations of the *Single-process Basic cache*: it manages to pipeline the *core* process with an II equal to 1, even in case of write-only accesses or long lines, since the AXI interfacing resides in the separate *memory interface* process.

The latency of the response to a hitting request depends on the *core* process only, therefore with this solution the best performance is achieved in case of write-only caches too.

In the case of caches which are accessed both in read and in write mode, it has not been possible to achieve an II of 1, due to dependencies on the cache memory. Given that a read-write cache implies at least one read access and one write access,

Figure 3.2: *Multi-processes Basic cache* architecture.

3.2 Implementation

The *Basic cache* is implemented in the form of a `C++14` [4] class, compatible with *Vitis™ HLS 2021.1*. All the configurable parameters are set through class template arguments.

The cache class is logically split into two parts:

- *Internals*: cache functionalities.

- *Interface*: APIs for managing requests and responses from application side.

Internals and *Interface* communicate with each other through a *Port* (Table 3.1), in a *Master/Slave* fashion:

- *Interface* sends to *Internals* a *request* (operation, address and write data).

- *Internals* sends to *Interface* a *response* (read data), after executing the requested operation.

Content	Description	Direction
Operation	Read/Write	*Internals → Interface*
Address	Index to be accessed	*Internals → Interface*
Write data	Data to be written to memory	*Internals → Interface*
Read data	Data read from memory	*Internals ← Interface*

Table 3.1: Data exchanged through *Port*.

Process modeling   HLS is intended for synthesizing sequential Software code, therefore it has been necessary to develop a novel technique for modeling multiprocess designs.

The proposed model follows the *Master/Slave* paradigm:

1. *Master* sends a request to *Slave*.

2. *Slave* executes the requested operation and optionally sends a response to *Master*.

Slave must be modeled as an infinite loop which waits for requests from *Master* before executing its functionality, while *Master* can be modeled as standard sequential code (or it can be in turn a *Slave* of another *Master*).

The parallelism between *Master* and *Slave* is modeled differently depending on the compilation target:

- *SW simulation*: each process is mapped to a `std::thread`.

- *HW synthesis*: each process is a dataflow function, in a dataflow region with the `disable_start_propagation` option disabled (which allow each function to run in parallel, without waiting for the completion of previous ones).

The distinction between simulation and synthesis code can be performed through the "`#ifdef __SYNTHESIS__`" preprocessor directive.

The communication between the two processes is performed through a *port*, which contains data flowing from *Master* to *Slave* (request) and from *Slave* to *Master* (response). Request and response are mapped to one or more FIFOs which are written from the transmitter and read from the receiver. `hls::stream` class by *Vitis™ HLS* can be used as FIFO implementation.

3.2.1 Internals

The *Internals* implementation differs between the *Single-process* and the *Multi-processes* implementations:

- *Single-process Basic cache*: single process which implements all the cache functionalities.

- *Multi-processes Basic cache*:

    - *Core* process: same as *Single-process Basic cache* process, but it does not directly access the AXI bus: it issues requests to the *memory interface* process through FIFOs.

    - *Memory interface* process: it accesses the AXI bus as requested by the *core* process.

Single-process Basic cache, with respect to the *Multi-processes* one, requires lower resource usage and better performance, when it is possible to schedule its process with an II equal to 1 (read-only accesses with line not larger than the maximum AXI interface bitwidth): therefore it is automatically instantiated whenever it is convenient.

Dataflow checking

Alternatively executing the *Multi-processes* or the *Single-process* code with traditional `if` statements would generate errors during the synthesis, particularly in the *Dataflow check* step (which checks if each `hls::stream` has a single reader and a single writer): the compiler builds both branches of the `if` statements, independently of the fact that one of them is never executed.

The problem has been solved through a wrapper class, which conditionally includes a `hls::stream` object, exploiting the template specialization mechanism.

Arrays partitioning

Cache memory (which stores the actual data) must be accessed one line per clock cycle: it is mapped to a BRAM array cyclically partitioned with a factor equal to the number of lines.

Helper data (e.g. `tag`, `valid`, `dirty`...) is stored to completely partitioned arrays, mapped to registers, in order to avoid dependencies as much as possible and get the best performance.

AXI optimizations

To exploit the *Vitis™ HLS* support to automatic port widening and burst accesses to AXI interface, every access to external DRAM accesses a whole cache line. The accessed

addresses Least Significant Bits (LSBs) are explicitly set to 0 so that synthesizer can infer that they are aligned to the line size.

If the cache line is at most equal to the maximum AXI interface width, it is accessed in a single request, otherwise it is accessed in multiple burst requests.

Read After Write dependencies

In case of read-write caches, the *Core* process II increases to 3 due to RAW dependencies on the cache BRAM.

To mitigate this issue the *RAW cache* has been developed: it is a single-line cache which provides the functions:

- `get_line`: in case of hit, read the *RAW cache* line; in case of miss, read the cache line.

- `set_line`: write both the *RAW cache* line and the cache line.

Cache memory is always accessed through the *RAW cache* and the `set_line` function is called once per iteration at most: if a cache line has been written, it is impossible that it is read in the next iteration, since the RAW cache would hit and return its line. This allows to falsify the RAW dependency with distance 1 on the cache memory (by setting to `false` the RAW inter-iteration dependencies and to `true` the RAW inter-iteration dependencies with distance 1).

This solution allows to schedule the cache process with an II equal to 2. The *RAW cache* could be extended to a fully-associative cache complying with the FIFO replacement policy, allowing to falsify the RAW dependency with distance 2 and achieving an II of 1.

A read-write cache implies that it is accessed at least two times per iteration (once in read mode, once in write mode), therefore, due to the issues discussed in Subsection 3.2.2 it is not possible to fully exploit the cache pipelining. In this case the cache II does not have a relevant impact on effective performance: *RAW cache* could not provide real advantages and it has not been included in the final design, to keep it simpler.

3.2.2 Interface

Interface provides APIs for managing requests and responses between application and cache:

- `get`: send a read request and read the response.

- `set`: send a write request.

To improve user-friendliness, similarly to *Ma's cache*, the `operator[]` has been overloaded so that a cache object can be used as a traditional array (e.g. `val = cache[i]` calls `val = cache.get(i)` and `cache[i] = val` calls `cache.set(i, val)`).

Deadlock prevention

The HLS scheduler is not able to infer the dependency between the request writing (W) and the response reading (R) in the `get` function (i.e. it is not aware that first the request has to be written, then it is necessary to wait for the cache latency and finally the response has to be read).

For that reason the scheduler optimizes the logic by inserting both W and R into the same pipeline stage. This leads to a deadlock: R is blocked since it reads from an empty FIFO (it cannot contain the response yet) and it blocks the whole stage, including W, making R wait for the response to a request which cannot be sent.

The deadlock has been fixed by inserting a clock operation between W and R (calling `ap_wait`), which forces W and R to separate pipeline stages.

Cache pipeline exploiting

At the steady state, in case of hit, the cache can process one request per cycle, thanks to its optimal pipelining (i.e. II equal to 1).

HLS is not aware of the dependency and latency between request writing (W) and response read (R), so it schedules R just after W (Figure 3.3a): at runtime R_i, which should be executed in the cycle following W_i, stalls, since the cache response has a latency (and W_{i+1} stalls too, by consequence).

W_{i+1} is executed after waiting for the full latency of the cache (Figure 3.3b) and the final result is that cache never receives multiple requests in consecutive cycles, it never reaches the steady state and its throughput is the same as if it were not pipelined.

(a) Static schedule.   (b) Runtime.

Figure 3.3: Stalling schedule of request writing and response reading.

To mitigate this issue the `ap_wait` between request write and response read has been replaced with `ap_wait_n(LATENCY)`, where `LATENCY` is an integer value set through a template parameter. This forces the scheduler to insert `LATENCY` clock cycles between W and R (Figure 3.4a), so that at runtime stalls are avoided (in case of hit) and one request per cycle is sent to cache (Figure 3.4b).

`LATENCY` is not set to a constant because its optimal value highly depends on memory access pattern and cache configuration, and can be determined by means of design exploration.

This is a partial solution: the `ap_wait` forces all the subsequent operations to wait: when there are multiple calls to `get` per iteration (e.g. A and B), W_B has to wait

(a) Static schedule. (b) Runtime.

Figure 3.4: Optimal schedule of request writing and response reading.

`LATENCY` cycles after W_A before being scheduled (Figure 3.5a). This situation makes the application loop II to increase, since it must guarantee the order of accesses to FIFOs (i.e. $W_{A,i+1}$ cannot be executed before $W_{B,i}$).

To actually fix this problem (with the schedule shown in Figure 3.5b), a mechanism for informing the scheduler about dependencies and latency between specific operations is probably needed, but this is not available in *Vitis™ HLS 2021.1*.

(a) Achieved static schedule. (b) Optimal static schedule.

Figure 3.5: Static schedules in case of multiple accesses per iteration.

4 Multi-levels cache

The *Multi-levels cache* is aimed at improving performance by making the memory hierarchy deeper, adding a faster L1 cache memory on top of it. This alternative approach has been proposed to overcome the difficulties, to fully exploit the optimal pipeline of the *Basic cache*, due to the scheduler unawareness about the latency between request writing and response reading (as explained in Section 3.2).

4.1 Architecture

The *Multi-levels cache* introduces a L1 cache inlined in the application logic (Figure 4.1): the scheduler exactly knows the latency of each L1 cache operation and can build an application pipeline which stalls in case of L1 miss only.

In order not to fall into the same cluttering issues of *Ma's cache*, the L1 cache is kept as simple as possible:

- Mapping policy: direct-mapped.

- Consistency policy: write-through.

The write-through consistency policy discards any advantage for write accesses, but given that simplicity is a priority and read accesses are usually more frequent than writes, and they suffer the most from the scheduling issues which lead to the introduction of the L1 cache, this has been considered the best trade-off.

4.2 Implementation

The *Multi-level cache* has been implemented adding the L1 cache to the *Basic cache*. It is possible to configure the number of L1 cache lines through the L1_CACHE_LINES template parameter. When it is set to 0, the resulting architecture is equivalent to the *Basic cache*.

4.2.1 Internals

The only difference with respect to the *Basic cache* implementation is that the response to a read request does not send a single word, but a whole cache line (therefore the data FIFO flowing from cache to application has been widened accordingly).

Figure 4.1: *Multi-levels cache* architecture.

4.2.2 Interface

The L1 cache is contained in the *Interface*: the newly introduced `get_line` function receives an address A in input and it returns the line to which A belongs. In particular, it first checks if A hits in the L1 cache: if so it reads the data from the L1 cache, otherwise it issues the request to the L2 cache.

It is still possible to use the same *Basic cache* APIs, which have been updated to support the L1 cache:

- `get`: it calls the `get_line` function and then returns the requested word.

- `set`: it sets L1 cache line to dirty, if it hits, and it forwards the request to the L2 cache.

19

5 Multi-ports cache

The computational core of many algorithms consists in a loop, which HLS can optimize with two techniques: *Pipelining* and *Unrolling*.

The *Basic* and *Multi-levels* caches are suitable for *Pipelining* since they complete one access per clock cycle, at the steady state, in case of hit, however they are not suitable for *Unrolling*, since they do not support concurrent accesses.

The *Multi-ports cache* has been specifically designed for adding support to multiple **concurrent accesses** to the same cache memory, allowing to efficiently **unroll** application loops.

5.1   Architecture

The *Multi-ports cache* is characterized by multiple ports accessed in parallel (Figure 5.1).

Each port has dedicated logic for communicating with the shared L2 cache and an independent L1 cache.

Multiple independent ports allow **removing dependencies** between different accesses to the cache. This brings the advantage of achieving better performance, making it possible to schedule multiple requests at the same time, without increasing the application loop II, but it also brings the disadvantage of not guaranteeing the expected ordering between different accesses. To guarantee the correct functionality the *Multi-ports* architecture is compatible with **read-only** accesses.

5.2   Implementation

The *Multi-ports cache* has been implemented extending the *Multi-levels cache*.

It is possible to configure the number of ports through the `PORTS` template parameter. When it is set to 1, the resulting architecture is equivalent to the *Multi-levels cache*.

5.2.1   Internals

To avoid dependencies issues, whenever `PORTS` is greater than 1, the *Multi-process Internals* architecture is generated.

Figure 5.1: *Multi-ports cache* architecture.

The *Core* process has been modified to serve requests coming from all the ports by inserting an unrolled loop which iterates over all the ports. HLS guarantees all the dependencies on cache data structures, and the resulting II of the *Core* process is equal to PORTS.

5.2.2 Interface

FIFOs between *Core* and application and L1 cache have been replaced with arrays of FIFOs and L1 caches, completely partitioned, so that they are independent.

Each call to get_line (which is in turn called by get) is automatically associated with a specific port by means of a member variable holding the port index and is updated after each access.

FIFOs accesses scheduling

Ideally the request write (W) and the response read (R) should be scheduled in parallel in the same cycle (Figure 5.2b). Due to the scheduler limitations (described in Subsection 3.2) it is not possible to achieve such a schedule, since there is a forced clock cycle between W and R, which delays all the subsequent operations.

The resulting schedule (Figure 5.2a) is almost equivalent to the one achieved with the *Basic cache* in case of multiple accesses per iteration (Figure 3.5a), with the difference that request and response FIFOs are distinct, since they belong to separate ports, therefore the scheduler does not have to ensure dependencies between subsequent reads and writes and application loop II does not increase.

At the steady state, in case of hit, one W and R are executed per cycle, allowing to fully exploit the L2 cache pipeline.

(a) Achieved static schedule.

(b) Parallel static schedule.

Figure 5.2: Static schedules in case of 2-ports cache.

5.3   Limitations

In some particular situations (e.g. when cache is explicitly accessed multiple times per iteration) the simulation of the generated circuit enters a deadlock. The source of this problem can be probably found in the port indexing and to be fixed may require more control over the operations scheduling, which is not provided by *Vitis™ HLS 2021.1*.

6 Results

The proposed cache architecture has been embedded in multiple *Vitis HLS* kernels implementing different algorithms, to evaluate both the performance gain and the resource usage of different cache configurations.

Each algorithm has been selected for its memory intensiveness and for its specific memory access patterns.

6.1 Simulation environment

Kernels have been synthesized by the C Synthesis in *Vitis™ HLS 2021.1*, targeting the `xcvu9p-flgb2104-2-e` part, running at a clock frequency of $250MHz$.

Vitis™ 2021.1 provides two main kind of simulation:

- Hardware Emulation: accurate, but slow.

- C/RTL Co-Simulation: fast, but not very accurate (especially for what concerns the AXI interface model).

HW Emulation has been used for determining the delay of the AXI interface (which is around 4 clock cycles). The AXI latency has been accordingly set to 3, so that the synthesizer can better optimize the circuit and Co-Simulation results match HW emulation as much as possible.

Synthesizer	C Synthesis in *Vitis™ HLS 2021.1*
Simulator	C/RTL Co-Simulation in *Vitis™ HLS 2021.1*
Flow target	`vitis`
Part codename	`xcvu9p-flgb2104-2-e`
Clock period	*4ns*
AXI latency	3

Table 6.1: Simulation environment configuration.

6.1.1   Reference memory models

The results have been compared with the output of synthesis and simulation of same algorithms implemented with different data access mechanisms: *global memory* (performance lower bound) and *local memory* (performance higher bound).

Global memory

The algorithms access data directly from external DRAM through AXI interface: this is the straightforward but slowest solution, therefore it determines the performance lower bound.

Local memory

All the data required by algorithms is stored to local BRAMs: it determines the performance higher bound, but it is unfeasible in general, due to the limited amount of BRAMs.

  With this solution the kernel:

1. Moves all the input data from DRAM to BRAMs.

2. Performs all the computations accessing data to and from BRAMs.

3. Moves all the output data from BRAMs to DRAM.

  The execution time of DRAM accesses is not of interest, therefore it has been subtracted from reported results.

Ma's cache

Ma's cache was designed for *Vivado™ HLS 2016.2*: with some minor changes it is possible to synthesize it with *Vitis HLS 2021.1*, but it would need some more optimizations to achieve the original performance in the new environment.

  The results reported in Ma's paper "Acceleration by Inline Cache for Memory-Intensive Algorithms on FPGA via High-Level Synthesis" and PhD thesis "Low power and high performance heterogeneous computing on FPGAs" are not comparable with the ones obtained in *Vitis HLS 2021.1*: the execution times of same algorithms with same configurations (i.e. global and local memory) differ up to one order of magnitude (most probably due to different AXI latency values), therefore also cache execution times would not be reliable for making comparisons.

  The lack of comparable results and the impossibility to generate new ones prevented from directly comparing the proposed cache results with the Ma's ones.

6.1.2 Collected data

The most relevant collected data concerns:

- Performance: evaluated in terms of execution time (i.e. the time at which the simulation of the algorithm terminates).

- Resource usage: evaluated in terms of number of used BRAMs, Digital Signal Processors (DSPs), Lookup Tables (LUTs) and Flip-Flops (FFs).

These values are approximate, since they come from C/RTL Co-Simulation and from the estimations performed by the C Synthesis, but in any case they can be meaningful for identifying some trends.

6.2 Matrix multiplication

The standard row-by-column *Matrix multiplication* algorithm (Algorithm 1) includes two memory access patterns: by rows (A and C) and by columns (B).

Each row of A matrix is accessed P times and then it is not accessed anymore: the most convenient A cache is composed of a single line which fits a matrix row, which is filled each time a new row is accessed and it hits until the next row is accessed.

Each column of B matrix is accessed P times: the B cache, to get a hit ratio greater than 0 needs to contain at least M lines and comply with the fully-associative mapping policy. The results reported by Ma used a direct-mapped cache with M lines each one containing P elements (so that it is as big as the B matrix).

C elements are accessed sequentially and only once: any single-line cache with n words per line would have a hit ratio of $\frac{n-1}{n}$.

The implementation used during the tests applies both pipelining and unrolling (with factor equal to the number of ports) to the innermost loop.

Algorithm 1 *Matrix multiplication* algorithm.

Require: $A \in \mathbb{R}^{N \times M}, B \in \mathbb{R}^{M \times P}, C \in \mathbb{R}^{N \times P}$
Ensure: $C = A \times B$
  procedure MULTIPLY(A, B, C)
    for $i = 0, \ldots, N - 1$ **do**
      for $j = 0, \ldots, P - 1$ **do**
        $tmp \leftarrow 0$
        for $k = 0, \ldots, M - 1$ **do**
          $tmp \leftarrow tmp + A[i][k] \cdot B[k][j]$
        end for
        $C[i][j] \leftarrow tmp$
      end for
    end for
  end procedure

6.2.1 16x16 matrices

In the case of *Matrix multiplication 16x16*, matrices A, B and C are sized 16×16 ($N = 16, M = 16, P = 16$).

This problem has been explored in two configurations: *Single-level cache configuration* (L2 caches only), and *Multi-level cache configuration* (L2 and L1 caches).

Single-level cache configuration

The cache sizes have been fixed with the values shown in Table 6.2. The `get` latency and the number of ports have been determined through design space exploration.

Matrix	Sets	Ways	Words per line	L1 lines	Hit ratio
A	1	1	16	0	99.6 %
B	16	1	16	0	99.6 %
C	1	1	16	0	93.8 %

Table 6.2: Single-level cache configuration for *Matrix multiplication 16x16*.

Figure 6.1 shows the execution time with respect to the `get` latency, for different numbers of ports.

It is worth noting that the `get` latency has a big impact on effective performance, especially in the single-port case (one order of magnitude). This makes clear that the cache process itself can run at high speed and the bottleneck is the scheduling of the FIFOs accesses.

Increasing the number of ports can provide significant advantages when the `get` latency is not optimal, because multi-port allow to schedule some cache requests in consecutive clock cycles.

Figure 6.1: Design space of *Matrix multiplication 16x16* (single-level).

The best performance is achieved by the single-port, since in this case the caches *core* process has an II of 1: with a `get` latency of 1 it is not possible to take full advantage of the `core` pipelining (as explained in Subsection 3.2.1), therefore the design keeps stalling even at the steady state (Figure 6.2a: a new request is written every multiple cycles) but the optimal `get` latency allows to fully exploit the pipelining and at every cycle one request is written and a new response is read (Figure 6.2b).

(a) Sub-optimal `get` latency of 1.

(b) Optimal `get` latency of 15.

Figure 6.2: Request and response waveforms for *Matrix multiplication 16x16* single-level and single-port.

Table 6.3 reports the data for the single-level cache configuration of different port numbers, each one set to its optimal `get` latency value (15 for the 1-port, 3 for the 2-ports and 2 for the 4-ports).

The single-port configuration fully exploits the underlying single L2 cache, therefore adding more ports not only increase the resource usage, but it also reduces performance since the cache II increases.

It is not clear why the estimated required BRAMs in the 2-ports case is much higher than the other cases.

	1-port	**2-ports**	**4-ports**
Execution time [ns]	17438	38570	49694
BRAM	90	165	90
DSP	3	6	12
LUT	57653	87437	118434
FF	26597	37686	39352

Table 6.3: Performance and resource usage of *Matrix multiplication 16x16* (single-level).

Multi-levels cache configuration

The cache sizes have been fixed with the values shown in Table 6.4. The `get` latency and the number of ports have been determined through design space exploration.

Matrix	Sets	Ways	Words per line	L1 lines	Hit ratio
A	1	1	16	1	99.6 %
B	1	1	16	16	99.6 %
C	1	1	16	0	93.8 %

Table 6.4: Multi-levels cache configuration for *Matrix multiplication 16x16*.

From Figure 6.3 it is clear that the `get` latency is not relevant is this case, since all the cache hits are on the L1 cache.

Increasing the number of ports allows to significantly improve performance, since the multiple L1 caches can run effectively in parallel. The higher is the number of ports, the lower is the hit ratio of each L1 cache: indefinitely increasing the number of ports is not always convenient; in this case 4 ports is the optimal configuration in terms of performance. More details about some simulated configurations are reported in Table 6.5.

Figure 6.3: Design space of *Matrix multiplication 16x16* (multi-levels).

	1-port	2-ports	4-ports	8-ports
get latency	8	7	7	1
Execution time [*ns*]	18986	10166	6458	7358
BRAM	129	165	237	381
DSP	3	6	12	24
LUT	58138	81779	118794	198678
FF	41961	101315	238374	581204

Table 6.5: Performance and resource usage of *Matrix multiplication 16x16* (multi-levels).

Summary

Table 6.6 reports most relevant figures about performance and resource usage of *Matrix multiplication 16x16*.

   The Proposed cache data is referred to the most performant cases of the single-level variant (with single port and `get` latency equal to 15) and of the multi-levels variant (with 4 ports and `get` latency equal to 7).

   The cost in terms of resource usage of the proposed cache is significant in terms of resource usage, particularly the number of LUTs and FFs is one order of magnitude more for the single-level and single-port configuration and two orders of magnitude for the multi-levels and multi-ports configuration, with respect to global memory. The single-level configuration allows reaching performance on par with the local memory and the multi-level configuration is more than two times faster, thanks to the unrolling.

	Global memory	Local memory	Proposed cache (single-level)	Proposed cache (multi-levels)
Execution time [*ns*]	30182	16916	17438	6458
BRAM	34	90	90	237
DSP	3	3	3	12
LUT	4421	26403	57653	118794
FF	4736	8829	26597	238374

Table 6.6: Performance and resource usage of *Matrix multiplication 16x16*.

6.2.2 32x32 matrices

To check whether the results scale with the problem size, in the case of *Matrix multiplication 32x32*, matrices *A*, *B* and *C* have been sized 32×32 ($N = 32, M = 32, P = 32$).

Single-level cache configuration

The cache sizes have been fixed with the values shown in Table 6.7. The `get` latency and the number of ports have been determined through design space exploration.

Matrix	Sets	Ways	Words per line	L1 lines	Hit ratio
A	1	1	32	0	99.9
B	32	1	32	0	99.9
C	1	1	32	0	96.9

Table 6.7: Single-level cache configuration for *Matrix multiplication 32x32*.

From Figure 6.4 it is possible to infer that the shapes of the Execution time - `get` latency plot of *Matrix multiplication 32x32* are equivalent to the ones obtained in the *16x16* case.

Figure 6.4: Design space of *Matrix multiplication 32x32* (single-level).

Results reported by Table 6.8 have been obtained by setting the `get` latency to value which provided the best performance (12 for the single-port cache and 3 for the dual-

port cache). The performance gain obtained by doubling the number of ports is not very significant.

	1-port	2-ports
Execution time [ns]	265226	229190
BRAM	90	90
DSP	3	6
LUT	126411	175727
FF	39871	47648

Table 6.8: Performance and resource usage of *Matrix multiplication 32x32* (single-level).

Multi-levels cache configuration

The cache sizes have been fixed with the values shown in Table 6.9. Since most of the read accesses are performed to L1 caches (Table 6.10), the `get` latency value does not have a big impact on the performance, therefore it has been fixed to 1.

The number of ports has been kept free for design space exploration.

Matrix	Sets	Ways	Words per line	L1 lines	get latency
A	1	1	32	1	1
B	1	1	32	32	1
C	1	1	32	0	1

Table 6.9: Multi-levels cache configuration for *Matrix multiplication 32x32*.

Matrix	L2 hit ratio [%]	L1 hit ratio [%]
A	0.7	99.2
B	0	99.9
C	96.9	0

Table 6.10: Multi-levels cache *Matrix multiplication 32x32* hit ratios.

	1-port	2-ports	4-ports	8-ports
Execution time [ns]	138974	71618	39398	30362
BRAM	90	90	90	90
DSP	3	6	12	24
LUT	113957	140117	183947	291738
FF	51004	75734	140269	336003

Table 6.11: Performance and resource usage of *Matrix multiplication 32x32* (multi-levels).

Summary

Table 6.12 compares the results obtained by different memory access mechanisms. The column *Cache* is referred to the most performant tested configuration: the multi-levels version with 8 ports, therefore the *Matrix multiplication* inner loop is unrolled by a factor 8. To make the higher bound comparable, there are two versions of the *Local memory* report: the first is referred to the case in which the *Matrix multiplication* inner loop is kept rolled, while in the second case the loop is unrolled with a factor 8 so that it is comparable with cache results.

	Global memory	Local memory	Local memory (unrolled)	Cache
Execution time [ns]	389498	131580	16920	30362
BRAM	34	90	90	90
DSP	3	3	24	24
LUT	4106	30589	121971	291738
FF	4699	11417	27866	336003

Table 6.12: Performance and resource usage of *Matrix multiplication 32x32*.

The cache can provide a great performance gain: it is one order of magnitude faster than the global memory and it is almost on par with the ideal case of the unrolled local memory. From Figure 6.5 it is possible to recognize the typical Pareto curve shape, which makes easy to find the desired trade-off: the more resources are employed, the more performance is delivered.
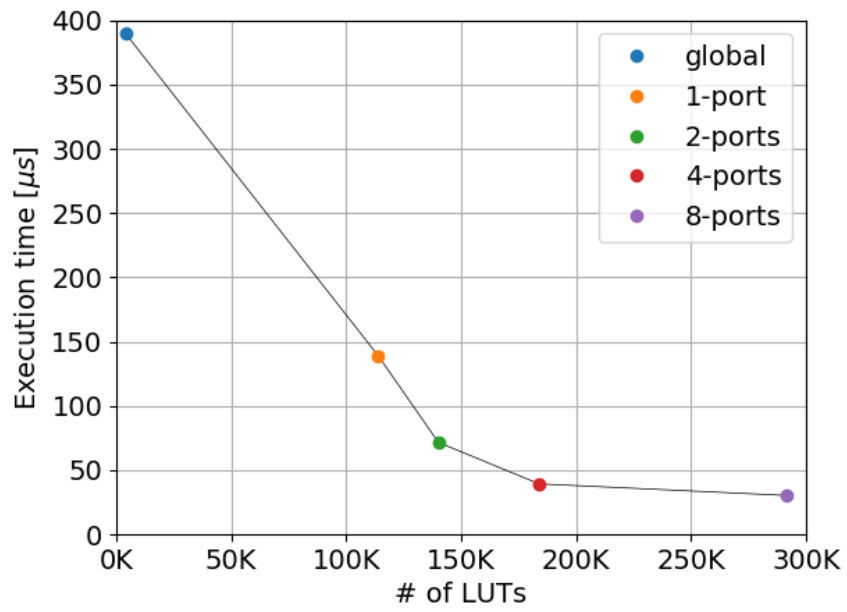
Figure 6.5: Pareto curve of *Matrix multiplication 32x32*.

6.3 Bitonic sorting

Bitonic sorting (Algorithm 2) is a sorting algorithm characterized by a high degree of parallelism, therefore it is suitable for Hardware implementation.

Algorithm 2 *Bitonic sorting* algorithm.

Require: $a \in \mathbb{R}^N, N = 2^n$; dir: sorting direction
Ensure: $a[i] \geq a[j], \forall i \geq j \wedge dir = true \vee a[i] \leq a[j], \forall i \geq j \wedge dir = false$
  procedure SORT(a, dir)
    for $b = 1, \ldots, n$ **do**
      for $s = i - 1, \ldots, 0$ **do**
        for $i = 0, \ldots, N/2 - 1$ **do**
          $dir_0 \leftarrow (i/2^{b-1})\&1$
          $dir_0 \leftarrow dir_0 | dir$
          $step \leftarrow 2^s$
          $pos \leftarrow 2i - (i\&(s - 1))$
          $a_0 \leftarrow a[pos]$
          $a_1 \leftarrow a[pos + step]$
          if $a_0 > a_1 \neq dir_0$ **then**
            $tmp \leftarrow a_0$
            $a_0 \leftarrow a_1$
            $a_1 \leftarrow tmp$
          end if
          $a[pos] \leftarrow a_0$
          $a[pos + step] \leftarrow a_1$
        end for
      end for
    end for
  end procedure

From the memory accesses point of view, each inner loop iteration:

1. $a[pos]$ is read.

2. $a[pos + step]$ is read.

3. $a[pos]$ is written.

4. $a[pos + step]$ is written.

The cache associated with a array should be set-associative with at least 2 sets, so that the interleaved accesses to pos and $pos + step$ do not overwrite the related cache lines.

In the design under test the inner loop have been pipelined, but, due to the data dependencies on a array, the pipeline performance is limited (i.e. II is greater than 1).

Due to the multiple accesses per iteration, it is not possible to exploit:

- Multiple ports: setting the number of ports to a number greater than one would result in a deadlock in C/RTL Co-Simulation, due to scheduling issues).

- `get` latency tuning: due to dependencies on a, each request to L1 cache is written only when the previous request is read.

For these reasons the `get` latency has been fixed to 2 (not 1, because it would lead the synthesizer to hang for unknown reasons) and the number of ports to 1.

6.3.1 128 elements

In the case of *Bitonic sorting 128*, array a contains 128 elements to be sorted ($n = 7$). The *Single-level cache configuration* is aimed at checking the performance of the L2 cache alone, while the *Multi-levels cache configuration* has the purpose to evaluate the boost given by adding a L1 cache on top of the L2.

Single-level cache configuration

For the *Single-level cache configuration* the cache associated with a is a fully-associative cache with two sets, so that the accesses to $pos + step$ index do not interfere with accesses to the pos index. The number of words per cache line has been kept free for design space exploration.
Table 6.13 summarizes the cache configuration.

Sets	Ways	L1 lines	get latency	Ports
1	2	0	2	1

Table 6.13: Single-level cache configuration for *Bitonic sorting 128*.

Table 6.14 shows the achieved results, with different number of words per cache line. Doubling the number of words approximately doubles the resource usage, while the performance grows at a much lower rate.

	8 words	16 words	32 words
Hit ratio [%]	93.8	96.9	98.4
Execution time [μs]	232	202	188
BRAM	16	30	30
LUT	20294	45481	91656
FF	6766	14403	24851

Table 6.14: Performance and resource usage of *Bitonic sorting 128* (single-level).

Multi-levels cache configuration

The *Multi-levels cache configuration* (Table 6.15) matches the *Single-level cache configuration*, with the only difference that a single-line L1 cache has been added on top of the memory hierarchy.

Sets	Ways	L1 lines	get latency	Ports
1	2	1	2	1

Table 6.15: Multi-levels cache configuration for *Bitonic sorting 128*.

Table 6.16 shows some results obtained with the *Multi-levels cache configuration*. Comparing these data with the *Single-level cache configuration* (Table 6.14) ones it is clear that the insertion of the L1 cache gives a further boost to performance, without increasing much the resource usage.

Execution time of the 16 words case is not reported because the Co-Simulation deadlocks with that specific configuration; its value would be probably between 150 μs and 170 μs.

	8 words	16 words	32 words
L1 hit ratio [%]	16.1	19.6	22.3
L2 hit ratio [%]	77.7	77.2	76.1
Execution time [μs]	194	-	130
BRAM	16	30	30
LUT	20745	46246	92740
FF	8082	18025	33066

Table 6.16: Performance and resource usage of *Bitonic sorting 128* (multi-levels).

Summary

Table 6.17 compares the results achieved with the most performant cache configuration (i.e. multi-levels with 32 words per line), with the higher and lower bounds for performance.

Cache is almost two times faster than the global memory, but it's still one order of magnitude slower than the local memory, and the resource usage is not negligible.

	Global memory	Local memory	Proposed cache
Execution time $[\mu s]$	244	23	130
BRAM	2	30	30
LUT	1743	3540	92740
FF	1088	3464	33066

Table 6.17: Performance and resource usage of *Bitonic sorting 128*.

From Figure 6.6 it is easy to conclude that the best trade-offs (between area and performance) belonging to the Pareto curve are the ones which include the L1 cache.

Figure 6.6: Pareto curve of *Bitonic sorting 128*.

6.4 2D convolution

Algorithm 3 is a possible implementation of *2D convolution* [1].

The algorithm accesses three matrices:

- *A*: read according to a sliding window pattern.

- *kernel*: read sequentially.

- *B*: written sequentially.

Accesses to *kernel* and *B* can be optimized by single-line caches, while *A* cache requires a more complex cache to get a sufficiently high hit ratio.

In the Hardware implementation the innermost loop have been pipelined, after moving the assignment to *B* in its latest iteration, to make the loop nest *perfect* (i.e. a loop contains only another loop without any external logic, reducing the HLS effort for pipelining).

Algorithm 3 *2D convolution* algorithm.

Require: $A \in \mathbb{R}^{N \times M}, kernel \in \mathbb{R}^{P \times Q}$
Ensure: $B \in \mathbb{R}^{N \times M} : B = A * kernel$
  procedure Conv($A, kernel$)
    for $i = 0, \ldots, N - 1$ **do**
      for $j = 0, \ldots, M - 1$ **do**
        $tmp \leftarrow 0$
        for $m = 0, \ldots, P - 1$ **do**
          for $n = 0, \ldots, Q - 1$ **do**
            $ii \leftarrow i + (Q/2 - m)$
            $jj \leftarrow j + (P/2 - n)$
            if $ii \geq 0$ & $ii < N$ & $jj \geq 0$ & $jj < M$ **then**
              $tmp \leftarrow tmp + A[ii][jj] \cdot kernel[m][n]$
            end if
          end for
        end for
        $B[i][j] \leftarrow tmp$
      end for
    end for
  end procedure

6.4.1 32x32 matrix and 3x3 kernel

The design under test has been set such that $A, B \in \mathbb{N}^{32 \times 32}$ and the $kernel \in \mathbb{N}^{3 \times 3}$.

The cache associated with A matrix is fully associative with 4 sets, to ensure that consecutive cache accesses do not overwrite previously loaded line. The other parameters have been left free for design space exploration.

kernel cache has been set to work as a line buffer: it contains a single line which fits the whole kernel, so that the first request load the kernel to cache and all the subsequent accesses are L1 cache hits.

B cache has the purpose of gathering multiple write requests.

The full configuration of caches is shown in Table 6.18.

Matrix	Sets	Ways	Words per line	L1 lines	Hit ratio [%]
A	1	4	Free	Free	?
kernel	1	1	16	1	100
B	1	1	32	0	96.9

Table 6.18: *kernel* and *B* caches configuration for *2D convolution*.

Single-level cache configuration

With the *Single-level cache configuration* the *A* cache is configured to have no L1 cache, while the number of words per line is kept free. The `get` latency has been set to 9 for both *A* and *kernel* caches.

Table 6.19 reports the collected results. The execution time with 32 words is missing since this configuration causes a deadlock in Co-Simulation.

	8 words	16 words	32 words
A **hit ratio** [%]	89.6	95.8	99.6
Execution time [μs]	45	40	-
BRAM	146	174	174
DSP	3	3	3
LUT	85574	91190	102887
FF	33803	37779	44578

Table 6.19: Performance and resource usage of *2D convolution* (single-level).

Multi-levels cache configuration

The *Multi-levels cache configuration* is equivalent to the *Single-level*, but a single-line L1 cache has been added.

The performance advantage with respect to the *Single-level configuration* is negligible, therefore it is not convenient to invest in more resources in this case.

	8 words	16 words	32 words
A **L1 hit ratio** [%]	59.6	63.8	66.0
A **L2 hit ratio** [%]	30.1	32.0	33.7
Execution time [μs]	43	40	39
BRAM	146	174	174
DSP	3	3	3
LUT	85640	91274	102970
FF	40288	50146	68720

Table 6.20: Performance and resource usage of *2D convolution* (multi-levels).

Summary

Table 6.21 compares the most performant cache configuration (i.e. multi-levels with 32 words per line) with the performance bounds: the achieved performance is almost equal to the performance higher bound.

	Global memory	Local memory	Cache
Execution time [μs]	66	37	39
BRAM	12	178	174
DSP	3	3	3
LUT	3504	30897	102970
FF	3257	10672	68720

Table 6.21: Performance and resource usage of *Convolution*.

The plot shown in Figure 6.7 suggests that the single-level 16 words per line configuration may be the most convenient in this case, since it offers almost the same performance and requires almost half of the FFs with respect to the multi-levels 32 words configuration.

Figure 6.7: Pareto curve of *2D convolution*.

7 Conclusions

This thesis work proven the possibility to implement multi-process designs in HLS.

The cache process itself can provide high performance thanks to its optimal pipelining (i.e. II of 1), but some issues arise from the proposed Inter-Process Communication protocol: *Vitis HLS 2021.1* do not provide any mechanism to inform the scheduler about the dependencies and latency between the request writing to a FIFO and the response reading from another FIFO, therefore the accesses to these data structures are not optimally scheduled by HLS, resulting in performance loss and in some cases even deadlocks.

The proposed workaround of forcing clock cycles to increase the latency between the FIFO write and the FIFO read allows to achieve optimal results in some specific situations, but it can not always work.

The collected results prove that in general the resource usage required by the proposed cache module is paid off in terms of performance, without requiring high design efforts for optimizing memory interfacing: it is enough to include the cache in the kernel and setup the parameters according to the desired performance and resource usage.

In some cases it is possible to achieve very high performance gain (e.g. *Matrix multiplication* achieved a reduction of execution time of more than one order of magnitude).

To overcome the limitations posed by the HLS tool it may be worth implementing the interface with the cache or the whole architecture at RTL, which provides full control on the operations scheduling.

The cache performance could be further improved by implementing a pre-fetching mechanism which analyses the memory access patterns and tries to load in advance the data, before it is requested.

A Source code

The source code produced during this thesis work is fully open source and can be found in the `git` repository hosted to https://github.com/brigio345/hls_cache.

The `src` directory contains the `cache` class and the related dependencies.

The `test` directory contains some test programs, including those used to collect the reported results; each of them can be synthesized and simulated in *Vitis HLS 2021.1* by executing the related script stored to `scripts`.

A.1 Cache

Listing A.1: Source code of `src/cache.h`

```
1  #ifndef CACHE_H
2  #define CACHE_H
3
4  /**
5   * \file   cache.h
6   *
7   * \brief   Cache module compatible with Vitis HLS 2021.1.
8   *
9   *    Cache module whose characteristics are:
10  *      - address mapping: set-associative;
11  *      - replacement policy: least-recently-used or
12  *          last-in first-out;
13  *      - write policy: write-back.
14  *
15  *    Advanced features:
16  *      - Multi-levels: L1 cache (direct-mapped, write-through).
17  *      - Multi-ports (read-only).
18  */
19
20 #include "address.h"
21 #include "replacer.h"
22 #include "l1_cache.h"
23 #define HLS_STREAM_THREAD_SAFE
24 #include "hls_stream.h"
25 #include "stream_cond.h"
26 #include "ap_utils.h"
27 #include "ap_int.h"
```

```
28  #include "utils.h"
29  #ifdef __SYNTHESIS__
30  #include "hls_vector.h"
31  #else
32  #include <thread>
33  #include <array>
34  #include <cassert>
35  #endif /* __SYNTHESIS__ */
36
37  #define MAX_AXI_BITWIDTH 512
38
39  template <typename T, bool RD_ENABLED, bool WR_ENABLED, size_t PORTS,
40      size_t MAIN_SIZE, size_t N_SETS, size_t N_WAYS,
41      size_t N_WORDS_PER_LINE, bool LRU,
42      size_t L1_CACHE_LINES, size_t LATENCY>
43  class cache {
44    private:
45      static const bool MEM_IF_PROCESS = (WR_ENABLED || (PORTS > 1) ||
46          ((sizeof(T) * N_WORDS_PER_LINE * 8) > MAX_AXI_BITWIDTH));
47      static const bool L1_CACHE = (L1_CACHE_LINES > 0);
48      static const size_t ADDR_SIZE = utils::log2_ceil(MAIN_SIZE);
49      static const size_t SET_SIZE = utils::log2_ceil(N_SETS);
50      static const size_t OFF_SIZE = utils::log2_ceil(N_WORDS_PER_LINE);
51      static const size_t TAG_SIZE = (ADDR_SIZE - (SET_SIZE + OFF_SIZE));
52      static const size_t WAY_SIZE = utils::log2_ceil(N_WAYS);
53
54      static_assert((RD_ENABLED || WR_ENABLED),
55          "RD_ENABLED and/or WR_ENABLED must be true");
56      static_assert((PORTS > 0), "PORTS must be greater than 0");
57      static_assert((!(WR_ENABLED && (PORTS > 1))),
58          "PORTS must be equal to 1 when WR_ENABLED is true");
59      static_assert(((MAIN_SIZE > 0) && ((1 << ADDR_SIZE) == MAIN_SIZE)),
60          "MAIN_SIZE must be a power of 2 greater than 0");
61      static_assert(((N_SETS > 0) && ((1 << SET_SIZE) == N_SETS)),
62          "N_SETS must be a power of 2 greater than 0");
63      static_assert(((N_WAYS > 0) && ((1 << WAY_SIZE) == N_WAYS)),
64          "N_WAYS must be a power of 2 greater than 0");
65      static_assert(((N_WORDS_PER_LINE > 0) &&
66          ((1 << OFF_SIZE) == N_WORDS_PER_LINE)),
67          "N_WORDS_PER_LINE must be a power of 2 greater than 0");
68      static_assert((MAIN_SIZE >= (N_SETS * N_WAYS * N_WORDS_PER_LINE)),
69          "N_SETS and/or N_WAYS and/or N_WORDS_PER_LINE are too big \
70          for the specified MAIN_SIZE");
71
72  #ifdef __SYNTHESIS__
73      template <typename TYPE, size_t SIZE>
74        using array_type = hls::vector<TYPE, SIZE>;
75  #else
76      template <typename TYPE, size_t SIZE>
77        using array_type = std::array<TYPE, SIZE>;
78  #endif /* __SYNTHESIS__ */
79
80      typedef address<ADDR_SIZE, TAG_SIZE, SET_SIZE, WAY_SIZE>
```

```
81        address_type;
82      typedef array_type<T, N_WORDS_PER_LINE> line_type;
83      typedef l1_cache<T, MAIN_SIZE, L1_CACHE_LINES, N_WORDS_PER_LINE>
84        l1_cache_type;
85      typedef replacer<LRU, address_type, N_SETS, N_WAYS,
86        N_WORDS_PER_LINE> replacer_type;
87
88      typedef enum {
89        READ_OP,
90        WRITE_OP,
91        READ_WRITE_OP,
92        STOP_OP,
93        NOP_OP
94      } op_type;
95
96 #if (defined(PROFILE) && (!defined(__SYNTHESIS__)))
97      typedef enum {
98        MISS,
99        HIT,
100       L1_HIT
101     } hit_status_type;
102 #endif /* (defined(PROFILE) && (!defined(__SYNTHESIS__))) */
103
104     typedef struct {
105       op_type op;
106       ap_uint<ADDR_SIZE> load_addr;
107       ap_uint<ADDR_SIZE> write_back_addr;
108       line_type line;
109     } mem_req_type;
110
111     ap_uint<(TAG_SIZE > 0) ? TAG_SIZE : 1> m_tag[N_SETS * N_WAYS];
112     bool m_valid[N_SETS * N_WAYS];
113     bool m_dirty[N_SETS * N_WAYS];
114     T m_cache_mem[N_SETS * N_WAYS * N_WORDS_PER_LINE];
115     hls::stream<op_type, 4> m_core_req_op[PORTS];
116     hls::stream<ap_uint<ADDR_SIZE>, 4> m_core_req_addr[PORTS];
117     hls::stream<T, 4> m_core_req_data[PORTS];
118     hls::stream<line_type, 4> m_core_resp[PORTS];
119     stream_cond<mem_req_type, 2, MEM_IF_PROCESS> m_mem_req;
120     stream_cond<line_type, 2, MEM_IF_PROCESS> m_mem_resp;
121     l1_cache_type m_l1_cache_get[PORTS];
122     replacer_type m_replacer;
123     unsigned int m_core_port;
124 #if (defined(PROFILE) && (!defined(__SYNTHESIS__)))
125     hls::stream<hit_status_type> m_hit_status;
126     int m_n_reqs = 0;
127     int m_n_hits = 0;
128     int m_n_l1_hits = 0;
129 #endif /* (defined(PROFILE) && (!defined(__SYNTHESIS__))) */
130
131   public:
132     cache() {
133 #pragma HLS array_partition variable=m_tag complete dim=1
```

46

```
134 #pragma HLS array_partition variable=m_valid complete dim=1
135 #pragma HLS array_partition variable=m_dirty complete dim=1
136 #pragma HLS array_partition variable=m_cache_mem cyclic \
137       factor=N_WORDS_PER_LINE dim=1
138 #pragma HLS array_partition variable=m_core_req_op complete
139 #pragma HLS array_partition variable=m_core_req_addr complete
140 #pragma HLS array_partition variable=m_core_req_data complete
141 #pragma HLS array_partition variable=m_core_resp complete
142 #pragma HLS array_partition variable=m_l1_cache_get complete
143     }
144
145     /**
146      * \brief Initialize the cache.
147      *
148      * \note  Must be called before calling \ref run.
149      */
150     void init() {
151       m_core_port = 0;
152       if (L1_CACHE) {
153         for (auto port = 0; port < PORTS; port++)
154           m_l1_cache_get[port].init();
155       }
156     }
157
158     /**
159      * \brief   Start cache internal processes.
160      *
161      * \param[in] main_mem  The pointer to the main memory.
162      *
163      * \note    In case of synthesis this must be
164      *          called in a dataflow region with
165      *          disable_start_propagation option,
166      *          together with the function in which
167      *          cache is accessed.
168      *
169      * \note     In case of C simulation this must be
170      *          executed by a thread separated from the
171      *          thread in which cache is accessed.
172      */
173     void run(T * const main_mem) {
174 #pragma HLS inline
175 #ifdef __SYNTHESIS__
176       run_core(main_mem);
177       if (MEM_IF_PROCESS)
178         run_mem_if(main_mem);
179 #else
180       std::thread core_thd([&]{run_core(main_mem);});
181       if (MEM_IF_PROCESS) {
182         std::thread mem_if_thd([&]{
183             run_mem_if(main_mem);
184             });
185
186         mem_if_thd.join();
```

47

```
187        }
188
189        core_thd.join();
190 #endif /* __SYNTHESIS__ */
191    }
192
193    /**
194     * \brief Stop cache internal processes.
195     *
196     * \note  Must be called after the function in which cache
197     *    is accessed has completed.
198     */
199    void stop() {
200      for (auto port = 0; port < PORTS; port++)
201        m_core_req_op[port].write(STOP_OP);
202    }
203
204    /**
205     * \brief   Request to read a whole cache line.
206     *
207     * \param[in] addr_main The address in main memory belonging to
208     *      the cache line to be read.
209     * \param[out] line The buffer to store the read line.
210     */
211    void get_line(const ap_uint<ADDR_SIZE> addr_main, line_type &line) {
212 #pragma HLS inline
213 #ifndef __SYNTHESIS__
214      assert(addr_main < MAIN_SIZE);
215 #endif /* __SYNTHESIS__ */
216
217      const auto port = m_core_port;
218      m_core_port = ((m_core_port + 1) % PORTS);
219
220      // try to get line from L1 cache
221      const auto l1_hit = (L1_CACHE &&
222          m_l1_cache_get[port].hit(addr_main));
223
224      if (l1_hit) {
225        m_l1_cache_get[port].get_line(addr_main, line);
226 #ifndef __SYNTHESIS__
227        m_core_req_op[port].write(NOP_OP);
228 #endif /* __SYNTHESIS__ */
229      } else {
230        // send read request to cache
231        m_core_req_op[port].write(READ_OP);
232        m_core_req_addr[port].write(addr_main);
233        // force FIFO write and FIFO read to separate
234        // pipeline stages to avoid deadlock due to
235        // the blocking read
236        ap_wait_n(LATENCY);
237        // read response from cache
238        m_core_resp[port].read(line);
239
```

```
240          if (L1_CACHE) {
241            // store line to L1 cache
242            m_l1_cache_get[port].set_line(addr_main, line);
243          }
244        }
245
246 #if (defined(PROFILE) && (!defined(__SYNTHESIS__)))
247        update_profiling(l1_hit ? L1_HIT : m_hit_status.read());
248 #endif /* (defined(PROFILE) && (!defined(__SYNTHESIS__))) */
249      }
250
251      /**
252       * \brief   Request to read a data element.
253       *
254       * \param[in] addr_main The address in main memory referring to
255       *      the data element to be read.
256       *
257       * \return    The read data element.
258       */
259      T get(const ap_uint<ADDR_SIZE> addr_main) {
260 #pragma HLS inline
261        line_type line;
262
263        // get the whole cache line
264        get_line(addr_main, line);
265
266        // extract information from address
267        address_type addr(addr_main);
268
269        return line[addr.m_off];
270      }
271
272      /**
273       * \brief   Request to write a data element.
274       *
275       * \param[in] addr_main The address in main memory referring to
276       *      the data element to be written.
277       * \param[in] data  The data to be written.
278       */
279      void set(const ap_uint<ADDR_SIZE> addr_main, const T data) {
280 #pragma HLS inline
281 #ifndef __SYNTHESIS__
282        assert(addr_main < MAIN_SIZE);
283 #endif /* __SYNTHESIS__ */
284
285        if (L1_CACHE) {
286          // inform L1 caches about the writing
287          m_l1_cache_get[0].notify_write(addr_main);
288        }
289
290        // send write request to cache
291        m_core_req_op[0].write(WRITE_OP);
292        m_core_req_addr[0].write(addr_main);
```

49

```
293        m_core_req_data[0].write(data);
294 #if (defined(PROFILE) && (!defined(__SYNTHESIS__)))
295        update_profiling(m_hit_status.read());
296 #endif /* (defined(PROFILE) && (!defined(__SYNTHESIS__))) */
297     }
298
299 #if (defined(PROFILE) && (!defined(__SYNTHESIS__)))
300     int get_n_reqs() const {
301       return m_n_reqs;
302     }
303
304     int get_n_hits() const {
305       return m_n_hits;
306     }
307
308     int get_n_l1_hits() const {
309       return m_n_l1_hits;
310     }
311
312     double get_hit_ratio() const {
313       if (m_n_reqs > 0)
314         return ((m_n_hits + m_n_l1_hits) /
315             static_cast<double>(m_n_reqs));
316
317       return 0;
318     }
319 #endif /* (defined(PROFILE) && (!defined(__SYNTHESIS__))) */
320
321   private:
322     /**
323      * \brief    Infinite loop managing the cache access
324      *       requests (sent from the outside).
325      *
326      * \param[in] main_mem  The pointer to the main memory (ignored
327      *       if \ref MEM_IF_PROCESS is \c true).
328      *
329      * \note     The infinite loop must be stopped by
330      *       calling \ref stop (from the outside)
331      *       when all the accesses have been completed.
332      */
333     void run_core(T * const main_mem) {
334 #pragma HLS inline off
335         // invalidate all cache lines
336         for (auto line = 0; line < (N_SETS * N_WAYS); line++)
337           m_valid[line] = false;
338
339         m_replacer.init();
340
341 CORE_LOOP:      while (1) {
342 #pragma HLS pipeline II=PORTS
343 INNER_CORE_LOOP:     for (auto port = 0; port < PORTS; port++) {
344           op_type op;
345 #ifdef __SYNTHESIS__
```

```
346            // get request and
347            // make pipeline flushable (to avoid deadlock)
348            if (m_core_req_op[port].read_nb(op)) {
349 #else
350            // get request
351            m_core_req_op[port].read(op);
352 #endif /* __SYNTHESIS__ */
353
354            // exit the loop if request is "end-of-request"
355            if (op == STOP_OP)
356              goto core_end;
357
358 #ifndef __SYNTHESIS__
359            if (op == NOP_OP)
360              continue;
361 #endif /* __SYNTHESIS__ */
362
363            // check the request type
364            const auto read = ((RD_ENABLED && (op == READ_OP)) ||
365                (!WR_ENABLED));
366
367            // in case of write request, read data to be written
368            const auto addr_main = m_core_req_addr[port].read();
369            T data;
370            if (!read)
371              data = m_core_req_data[port].read();
372
373            // extract information from address
374            address_type addr(addr_main);
375
376            auto way = hit(addr);
377            const auto is_hit = (way != -1);
378
379            if (!is_hit)
380              way = m_replacer.get_way(addr);
381
382            addr.set_way(way);
383            m_replacer.notify_use(addr);
384
385            line_type line;
386            if (is_hit) {
387              // read from cache memory
388              get_line(m_cache_mem,
389                  addr.m_addr_cache,
390                  line);
391            } else {
392              // read from main memory
393              auto op = READ_OP;
394              // build write-back address
395              address_type write_back_addr(m_tag[addr.m_addr_line],
396                  addr.m_set,
397                  0, addr.m_way);
398              // check if write back is necessary
```

51

```
399              if (WR_ENABLED && m_valid[addr.m_addr_line] &&
400                  m_dirty[addr.m_addr_line]) {
401                // get the line to be written back
402                get_line(m_cache_mem,
403                    write_back_addr.m_addr_cache,
404                    line);
405
406                op = READ_WRITE_OP;
407              }
408
409              const mem_req_type req = {op, addr.m_addr_main,
410                  write_back_addr.m_addr_main, line};
411
412              if (MEM_IF_PROCESS) {
413                // send read request to
414                // memory interface and
415                // write request if
416                // write-back is necessary
417                m_mem_req.write(req);
418
419                // force FIFO write and
420                // FIFO read to separate
421                // pipeline stages to
422                // avoid deadlock due to
423                // the blocking read
424                ap_wait();
425
426                // read response from
427                // memory interface
428                m_mem_resp.read(line);
429              } else {
430                execute_mem_if_req(main_mem,
431                    req, line);
432              }
433
434              m_tag[addr.m_addr_line] = addr.m_tag;
435              m_valid[addr.m_addr_line] = true;
436              m_dirty[addr.m_addr_line] = false;
437
438              m_replacer.notify_insertion(addr);
439
440              if (read) {
441                // store loaded line to cache
442                set_line(m_cache_mem,
443                    addr.m_addr_cache,
444                    line);
445              }
446            }
447
448            if (read) {
449              // send the response to the read request
450              m_core_resp[port].write(line);
451            } else {
```

```
452              // modify the line
453              line[addr.m_off] = data;
454
455              // store the modified line to cache
456              set_line(m_cache_mem,
457                  addr.m_addr_cache,
458                  line);
459
460
461              m_dirty[addr.m_addr_line] = true;
462          }
463
464 #if (defined(PROFILE) && (!defined(__SYNTHESIS__)))
465          m_hit_status.write(is_hit ? HIT : MISS);
466 #endif /* (defined(PROFILE) && (!defined(__SYNTHESIS__))) */
467 #ifdef __SYNTHESIS__
468          }
469 #endif /* __SYNTHESIS__ */
470        }
471      }
472
473 core_end:
474      // synchronize main memory with cache memory
475      if (WR_ENABLED)
476        flush();
477
478      // make sure that flush has completed before stopping
479      // memory interface
480      ap_wait();
481
482      if (MEM_IF_PROCESS) {
483        // stop memory interface
484        line_type dummy;
485        m_mem_req.write({STOP_OP, 0, 0, dummy});
486      }
487    }
488
489    /**
490     * \brief   Infinite loop managing main memory
491     *      access requests (sent from \ref run_core).
492     *
493     * \param[in] main_mem  The pointer to the main memory.
494     *
495     * \note    \p main_mem must be associated with
496     *      a dedicated AXI port.
497     *
498     * \note    The infinite loop is stopped by
499     *      \ref run_core when it is in turn stopped
500     *      from the outside.
501     */
502    void run_mem_if(T * const main_mem) {
503 #pragma HLS inline off
504
```

53

```
505 MEM_IF_LOOP:      while (1) {
506 #pragma HLS pipeline off
507         mem_req_type req;
508         // get request
509         m_mem_req.read(req);
510
511         // exit the loop if request is "end-of-request"
512         if (req.op == STOP_OP)
513           break;
514
515         line_type line;
516         execute_mem_if_req(main_mem, req, line);
517
518         if ((req.op == READ_OP) || (req.op == READ_WRITE_OP)) {
519           // send the response to the read request
520           m_mem_resp.write(line);
521         }
522       }
523
524     }
525
526     /**
527      * \brief   Execute memory access(es) specified in
528      *        \p req.
529      *
530      * \param[in] main_mem  The pointer to the main memory.
531      * \param[in] req The request to be executed.
532      * \param[out] line The buffer to store the read line.
533      *
534      * \note     \p main_mem must be associated with
535      *        a dedicated AXI port.
536      */
537     void execute_mem_if_req(T * const main_mem,
538         const mem_req_type &req, line_type &line) {
539 #pragma HLS inline
540       if ((req.op == READ_OP) || (req.op == READ_WRITE_OP)) {
541         // read line from main memory
542         get_line(main_mem, req.load_addr, line);
543       }
544
545       if (WR_ENABLED && ((req.op == WRITE_OP) ||
546           (req.op == READ_WRITE_OP))) {
547         // write line to main memory
548         set_line(main_mem, req.write_back_addr, req.line);
549       }
550     }
551
552     /**
553      * \brief   Check if \p addr causes an HIT or a MISS.
554      *
555      * \param[in] addr  The address to be checked.
556      *
557      * \return    hitting way on HIT.
```

54

```
558        * \return     -1 on MISS.
559        */
560       inline int hit(const address_type &addr) const {
561 #pragma HLS inline
562           auto addr_tmp = addr;
563           auto hit_way = -1;
564           for (auto way = 0; way < N_WAYS; way++) {
565             addr_tmp.set_way(way);
566             if (m_valid[addr_tmp.m_addr_line] &&
567                 (addr_tmp.m_tag == m_tag[addr_tmp.m_addr_line])) {
568               hit_way = way;
569             }
570           }
571
572           return hit_way;
573       }
574
575       /**
576        * \brief Write back all valid dirty cache lines to main memory.
577        */
578       void flush() {
579 #pragma HLS inline
580           for (auto set = 0; set < N_SETS; set++) {
581             for (auto way = 0; way < N_WAYS; way++) {
582               const address_type addr(
583                   m_tag[set * N_WAYS + way],
584                   set, 0, way);
585               // check if line has to be written back
586               if (m_valid[addr.m_addr_line] &&
587                   m_dirty[addr.m_addr_line]) {
588                 // write line back
589                 line_type line;
590
591                 // read line
592                 get_line(m_cache_mem,
593                     addr.m_addr_cache,
594                     line);
595
596                 // send write request to memory
597                 // interface
598                 m_mem_req.write({WRITE_OP, 0,
599                     addr.m_addr_main,
600                     line});
601
602                 m_dirty[addr.m_addr_line] = false;
603               }
604             }
605           }
606       }
607
608       void get_line(const T * const mem,
609           const ap_uint<(ADDR_SIZE > 0) ? ADDR_SIZE : 1> addr,
610           line_type &line) {
```

55

```
611 #pragma HLS inline
612       const T * const mem_line = &(mem[addr & (-1U << OFF_SIZE)]);
613
614       for (auto off = 0; off < N_WORDS_PER_LINE; off++) {
615 #pragma HLS unroll
616          line[off] = mem_line[off];
617       }
618    }
619
620    void set_line(T * const mem,
621          const ap_uint<(ADDR_SIZE > 0) ? ADDR_SIZE : 1> addr,
622          const line_type &line) {
623 #pragma HLS inline
624       T * const mem_line = &(mem[addr & (-1U << OFF_SIZE)]);
625
626       for (auto off = 0; off < N_WORDS_PER_LINE; off++) {
627 #pragma HLS unroll
628          mem_line[off] = line[off];
629       }
630    }
631
632 #if (defined(PROFILE) && (!defined(__SYNTHESIS__)))
633    void update_profiling(const hit_status_type status) {
634      m_n_reqs++;
635
636      if (status == HIT)
637        m_n_hits++;
638      else if (status == L1_HIT)
639        m_n_l1_hits++;
640    }
641 #endif /* (defined(PROFILE) && (!defined(__SYNTHESIS__))) */
642
643    class square_bracket_proxy {
644      private:
645        cache *m_cache;
646        const ap_uint<ADDR_SIZE> m_addr_main;
647      public:
648        square_bracket_proxy(cache *c,
649            const ap_uint<ADDR_SIZE> addr_main):
650          m_cache(c), m_addr_main(addr_main) {
651 #pragma HLS inline
652          }
653
654        operator T() const {
655 #pragma HLS inline
656          return get();
657        }
658
659        square_bracket_proxy &operator=(const T data) {
660 #pragma HLS inline
661          set(data);
662          return *this;
663        }
```

```
664
665          square_bracket_proxy &operator=(
666              const square_bracket_proxy &proxy) {
667 #pragma HLS inline
668            set(proxy.get());
669            return *this;
670          }
671
672        private:
673          T get() const {
674 #pragma HLS inline
675            return m_cache->get(m_addr_main);
676          }
677
678          void set(const T data) {
679 #pragma HLS inline
680            m_cache->set(m_addr_main, data);
681          }
682        };
683
684    public:
685      square_bracket_proxy operator[](const ap_uint<ADDR_SIZE> addr_main) {
686 #pragma HLS inline
687        return square_bracket_proxy(this, addr_main);
688      }
689 };
690
691 #endif /* CACHE_H */
```

Bibliography

[1] Song Ho Ahn. *Convolution*. June 2018. URL: http://www.songho.ca/dsp/convolution/convolution.html#convolution_2d (visited on Oct. 12, 2021).

[2] Xilinx Inc. *Vitis High-Level Synthesis User Guide*. Mar. 2021. URL: https://www.xilinx.com/support/documentation/sw_manuals/xilinx2020_2/ug1399-vitis-hls.pdf.

[3] Xilinx Inc. *Vitis High-Level Synthesis User Guide*. Aug. 2021. URL: https://www.xilinx.com/support/documentation/sw_manuals/xilinx2021_1/ug1399-vitis-hls.pdf.

[4] ISO. *ISO/IEC 14882:2014 Information technology — Programming languages — C++*. Fourth. Dec. 2014.

[5] Liang Ma. "Low power and high performance heterogeneous computing on FPGAs". PhD thesis. Politecnico di Torino, Feb. 2019. DOI: http://dx.doi.org/10.6092%2Fpolito%2Fporto%2F2727228.

[6] Liang Ma, Luciano Lavagno, Mihai Lazarescu, and Arslan Arif. "Acceleration by Inline Cache for Memory-Intensive Algorithms on FPGA via High-Level Synthesis". In: *IEEE Access* PP (Sept. 2017), pp. 1–1. DOI: 10.1109/ACCESS.2017.2750923.